

Executive Summary

KEY TAKEAWAY

Harvey AI solves real legal workflow problems—particularly document review and contract analysis—but its complete lack of customization, single-vendor model dependency, and significant data portability gaps create a profile where short-term productivity gains come at the cost of organizational AI capacity building and strategic flexibility.

Harvey AI delivers genuine value for specific legal workflows—large-scale document review, contract analysis, and legal research—where its domain-tuned approach produces measurably useful outputs. That value is real and shouldn't be dismissed. However, the evaluation reveals a consistent pattern across nearly every criterion: Harvey trades organizational control, transparency, and long-term flexibility for a polished, opinionated experience. It scored 'No' on customization and only 'Partially' on every other criterion. The tool provides no ability to adjust system prompts, select models, modify temperature settings, or shape AI behavior in any meaningful way. For a Fortune 100 legal department investing in AI capability, this is a significant constraint—it means the organization is entirely dependent on Harvey's judgment about how AI should be applied to their work.

The strategic risks compound when you examine the full picture. Harvey's AI layer rests on a single-vendor dependency (OpenAI), meaning your organization inherits concentration risk it cannot manage or even fully observe. The data portability situation follows what the evaluation describes as a 'roach motel' pattern—standard formats go in, but there is no bulk export capability for the institutional knowledge, custom workflows, or annotated outputs your teams build inside the platform. Meanwhile, the transparency Harvey provides is selective: users see what context is being used but not the underlying prompts, model configurations, or architectural decisions that determine output quality. This opacity makes it difficult for legal teams to develop the deeper AI fluency—understanding why outputs succeed or fail—that will differentiate high-performing organizations over the next several years.

From a capacity-building perspective, Harvey builds basic AI interaction skills (prompt iteration, output evaluation, citation checking) but actively abstracts away the mechanical understanding that enables teams to work effectively with AI across tools and contexts. The learning is Harvey-specific, not transferable. This tool is best suited for organizations that have made a deliberate strategic decision to outsource their legal AI capability to a vendor and prioritize near-term productivity over long-term organizational learning. It is a poor fit for organizations that view AI fluency as a core competency to be developed internally, that require multi-model flexibility to manage vendor concentration risk, or that need confidence in their ability to switch providers as this market rapidly evolves.



Detailed Evaluation

Assessment across six critical criteria for AI vendor evaluation

SEE

Transparency & Visibility

~ Partially Met

ASSESSMENT

Harvey provides meaningful but selective transparency. Users get visibility into *what* context is used (Vault file/folder selection) and a process-level view of reasoning via "Thinking States," but the core AI mechanics are deliberately abstracted. System prompts are hidden, replaced by a curated visualization of reasoning steps. Model selection logic—including the ~100 parallel model calls and routing between custom fine-tuned models and general-purpose LLMs—is entirely opaque to the user, exposed only as a high-level "mode" toggle (draft vs. assist). Standard inference parameters (temperature, top_p) are not configurable, and RAG implementation details (chunking strategy, embedding models, vector database) are undisclosed. The citation system is genuinely strong, providing page-level source attribution with extracted text, which supports verifiability of outputs even if the underlying mechanics remain hidden. This is a deliberate architectural choice: Harvey is trading transparency for controlled quality in a high-stakes legal domain. [1, 2, 3, 4, 1, 5, 3, 6, 6, 7, 5, 8, 9, 10, 11, 12, 13, 14, 15, 16, 3, 3, 5, 17, 7, 7, 18, 19, 20, 21, 22, 23, 7, 5, 7, 24, 5, 25, 26, 27, 28, 29, 30, 31, 7, 32, 33, 34, 35]

TRADE-OFFS

● WHAT YOU GAIN

A tightly controlled, legally-optimized system where Harvey's domain expertise manages the complexity of model selection, inference tuning, and retrieval—reducing the risk of misconfiguration in high-stakes legal work. The citation engine and Thinking States provide output-level verifiability without requiring users to understand AI internals. For firms that want reliable legal AI without building internal AI engineering capacity, this abstraction is a feature.

● WHAT YOU GIVE UP

The organization cannot independently audit, benchmark, or understand the AI mechanics driving its legal work product. If Harvey silently changes models, adjusts inference parameters, or modifies retrieval strategies, the firm has no way to detect or evaluate those changes. This also impedes the organization's ability to develop deep AI fluency—lawyers learn to use Harvey's interface, not to understand how AI-assisted legal reasoning actually works. The ~100 parallel model calls per prompt also raise unanswered questions about cost structure and whether Harvey has incentives to reduce that compute intensity over time in ways that degrade quality invisibly.

QUESTIONS FOR VENDOR

1. Will Harvey provide enterprise customers with visibility into which specific models (including version numbers) are invoked for each task, and will you commit to notifying customers before changing model routing logic or swapping underlying models?
2. Can enterprise customers access or audit the system prompts, or at minimum receive documentation of the instructions and constraints applied to the model before their input is processed—particularly for high-stakes workflows like contract analysis or litigation research?
3. What specific technical details of your RAG implementation—embedding models, chunking strategies, retrieval scoring thresholds—can you share under NDA with enterprise customers who need to satisfy their own AI governance and audit requirements?

CHANGE

Customizability & Control

Not Met

ASSESSMENT

Harvey provides no meaningful customization of its AI layer. The evidence is clear and consistent across all four dimensions: users cannot adjust system prompts, select models, modify retrieval logic (chunk size, algorithm, k-value), or tune inference parameters (temperature, top_p). The API exposes only high-level functional toggles (``mode``, ``include_citations``, ``stream``). What Harvey calls "customization"—Workflow Builder sequencing, Vault file uploads, natural-language Playbooks—operates entirely above the AI layer, configuring *inputs* to a fully managed, opaque system rather than modifying how the AI itself behaves. The "agentic reasoning" approach to conditional logic means even workflow branching is delegated to the model rather than controlled by the user. Custom function calling is not yet generally available, and there is no client-accessible fine-tuning capability. ^[6],

3, 3, 2, 1, 6, 1, 33, 34, 35, 14, 4, 32, 36, 13, 3, 37, 12, 7, 38, 7, 5, 24, 11, 9, 1, 31, 7, 7, 5, 7, 25, 26, 30, 39, 31, 31, 18, 40, 41, 42, 5, 17, 5, 7, 5, 5, 43, 44, 45, 46, 47, 20]

TRADE-OFFS

● WHAT YOU GAIN

A fully managed, opaque AI layer eliminates the need for in-house ML engineering talent to operate the tool. Harvey's automated model orchestration ("100+ model calls" per query, mixture-of-experts routing) likely delivers strong out-of-the-box performance for standard legal tasks without any tuning burden. Legal professionals can use Playbooks and Workflows without technical expertise. The locked-down architecture also reduces the risk of misconfiguration by end users in a high-stakes legal context.

● WHAT YOU GIVE UP

The organization surrenders all control over AI behavior to Harvey. You cannot optimize for your firm's specific quality standards, risk tolerances, or cost sensitivities at the model level. If Harvey changes its underlying models, prompts, or retrieval logic, you have no visibility into or control over those changes. This creates deep dependency: your legal workflows are shaped entirely by Harvey's engineering decisions. It also impedes organizational AI learning—your teams build fluency in Harvey's interface, not in the underlying AI mechanics that will define competitive advantage. If Harvey's opaque orchestration produces suboptimal results for a specific use case, your only recourse is to rephrase your prompt or wait for Harvey to fix it.

QUESTIONS FOR VENDOR

1. Harvey's automated model orchestration selects models per sub-task—what visibility do customers have into which models were used for a given output, and what is the notification process when underlying models are changed or deprecated?
2. Is there a roadmap for exposing inference parameters (temperature, top_p) or system prompt customization to enterprise customers, particularly for use cases where firms need deterministic or highly constrained outputs?
3. When Harvey Agents with custom function calling reaches general availability, will customers be able to define their own tool schemas and connect to arbitrary external APIs, or will tool use remain limited to Harvey-defined platform tools?

ADAPT

Vendor Agility & Innovation

~ Partially Met

ASSESSMENT

Harvey demonstrates genuine agility in incorporating advances *within the OpenAI ecosystem*—the 12-hour Deep Research integration is impressive, and the Workflow Engine's modular architecture enables rapid feature deployment. However, Harvey's exclusive dependency on OpenAI is a significant constraint on true vendor agility. When "keeping pace with AI advancement" means adopting the best available model for a given

task—regardless of provider—Harvey structurally cannot do this. Anthropic's Claude, Google's Gemini, and open-source models like Llama are excluded entirely, meaning Harvey's agility is bounded by OpenAI's own release cadence and capability trajectory. The fine-tuning investment risk is real and documented: millions spent customizing a model that was then surpassed by a standard release. This is not a vendor that can pivot to the best available foundation model; it is a vendor that moves fast within a single provider's orbit. [48, 49, 7, 50, 7, 51, 52, 23, 53, 21, 54, 19, 45, 5, 55, 5, 56, 5, 20, 44, 38, 57, 2, 14, 7, 25, 5, 11, 12, 13, 5, 7, 58, 59, 9, 8, 15, 10, 28, 7, 4, 27, 60, 30, 16, 7, 24, 5, 7, 5, 61, 1, 5, 6, 29, 47, 5]

TRADE-OFFS

● WHAT YOU GAIN

A deeply integrated, enterprise-grade legal AI platform that rapidly incorporates OpenAI's latest capabilities (Deep Research, reasoning models) with legal-domain fine-tuning, modular workflow architecture, and a continuous release cadence—delivering frontier OpenAI features faster than most competitors can.

● WHAT YOU GIVE UP

Model optionality. If a non-OpenAI model becomes materially superior for legal reasoning, summarization, or any core task, Harvey cannot adopt it. The organization is making a second-order bet—not just on Harvey, but on OpenAI maintaining frontier leadership indefinitely. The fine-tuning sunk cost dynamic also creates perverse incentives for Harvey to stick with custom models even when newer base models outperform them.

QUESTIONS FOR VENDOR

1. What is your technical and contractual roadmap for supporting non-OpenAI foundation models (e.g., Anthropic Claude, Google Gemini, open-source models), and what structural barriers exist to multi-model support?
2. When a new OpenAI model release outperforms your existing fine-tuned model, what is your typical validation and migration timeline—and can you provide specific examples of how long past model transitions have taken from release to production deployment?
3. How do you evaluate whether your fine-tuned models still outperform newer base models, and what is your policy when they don't—do you deprecate the fine-tuned version or continue running it?

USE

Genuine Utility & Value

~ Partially Met

ASSESSMENT

Harvey demonstrates genuine utility for specific legal workflows—particularly large-scale document review, contract analysis, and legal research—with the Bridgewater case study (2 days ' 2 hours, 95%+ time savings) providing the strongest concrete evidence. However, the case for Harvey over increasingly capable commodity models remains contested. The 97% lawyer preference rate and 83% hallucination reduction figures are self-reported internal benchmarks, not independent evaluations, and multiple practitioners describe Harvey as

a "wrapper" whose differentiation erodes with each foundation model update. The real, defensible utility appears concentrated in two areas: (1) RAG-based integration with a firm's private document management systems (iManage, NetDocuments), which generic models cannot replicate without custom engineering, and (2) legal-specific fine-tuning that improves output formatting and domain accuracy. Outside these areas—particularly for single-document analysis—the marginal value over Claude or GPT-4 with good prompting is unproven. Adoption challenges are also a genuine concern: user reports of slow attorney adoption and a steep learning curve suggest that realized productivity gains may fall well short of theoretical maximums. [48, 62, 63, 7, 61, 5, 53, 64, 13, 5, 5, 21, 65, 1, 5, 54, 5, 56, 20, 23, 66, 7, 14, 9, 67, 12, 7, 68, 27, 30, 25, 7, 24, 11, 7, 69, 49, 70, 17, 71, 72, 73, 74, 5, 55, 75, 5, 45, 19]

TRADE-OFFS

● WHAT YOU GAIN

A purpose-built legal AI platform with enterprise-grade security, private document search (Vault/RAG), legal-specific fine-tuning that reduces hallucinations on case law, and pre-built workflows for high-value tasks like due diligence and contract review—all without requiring internal engineering to build and maintain a custom LLM stack.

● WHAT YOU GIVE UP

Premium pricing for capabilities that partially overlap with commodity foundation models improving rapidly; risk that the differentiation gap narrows with each GPT/Claude update; potential adoption friction that limits realized ROI; and organizational dependency on Harvey's platform rather than building internal AI fluency that could be applied more broadly across the business.

QUESTIONS FOR VENDOR

1. Can you provide independently verified (not self-reported) performance benchmarks comparing Harvey's outputs against current frontier models (GPT-4o, Claude 3.5 Sonnet) on identical legal tasks, including accuracy, hallucination rates, and time-to-completion?
2. What is the median time savings per user across your entire customer base—not the top decile—and what percentage of licensed seats are actively used on a weekly basis after 6 months of deployment?
3. Given that your fine-tuning advantage may erode as foundation models improve, what is your concrete product roadmap for maintaining differentiation—and how much of your current value-add comes from the model layer versus the workflow/integration layer?

LEARN

Transferable Knowledge

~ Partially Met

ASSESSMENT

Harvey builds basic AI interaction literacy—iterating on prompts, evaluating outputs, understanding citation grounding—but actively abstracts away the deeper technical principles that constitute durable AI fluency. Users learn to prompt within Harvey's interface and gain exposure to concepts like RAG and model orchestration through blog content, but they never touch the mechanics: they don't select models, configure retrieval pipelines,

manage context windows, or evaluate trade-offs between inference cost and quality. The API exists and uses standard REST patterns (Bearer auth, multipart form uploads), which provides some transferable technical skill for developer-facing teams, but the absence of public SDKs and the proprietary orchestration layer ("100+ model calls per prompt") mean even API users are learning Harvey's abstraction, not foundational AI engineering. Harvey Academy and Harvey Guide are explicitly designed to teach users how to operate Harvey, not how AI works—the educational content on model selection and latency trade-offs lives in engineering blog posts that most end users will never read. [4, 7, 14, 27, 16, 35, 28, 7, 26, 29, 76, 9, 7, 33, 11, 13, 77, 5, 12, 25, 63, 69, 24, 17, 5, 58, 61, 78, 47, 79, 43, 60, 36, 57, 2, 32, 37, 6, 31, 41, 41, 34, 1, 5, 30, 7, 7, 10, 8, 5, 5, 7, 44, 20, 5, 5]

TRADE-OFFS

● WHAT YOU GAIN

Lawyers and professional staff become productive with AI-assisted legal work quickly, without needing to understand underlying AI mechanics. The abstraction layer removes friction and reduces error rates for non-technical users, delivering immediate workflow value.

● WHAT YOU GIVE UP

The organization builds proficiency in Harvey's interface rather than in AI itself. If Harvey is displaced or the firm wants to build internal AI capabilities, users will lack the foundational skills—prompt architecture, RAG design, model evaluation, context management—that would transfer. The firm becomes dependent on Harvey's engineering team for the AI competence that increasingly matters strategically.

QUESTIONS FOR VENDOR

1. Does Harvey plan to expose any of its model selection, orchestration logic, or RAG configuration to end users—even in an 'advanced mode'—so power users can develop deeper AI fluency rather than only consuming abstracted outputs?
2. Beyond Harvey Academy's tool-specific training, does Harvey offer or plan to offer educational content that explicitly teaches generalizable AI principles (e.g., how context windows work, when to use different model architectures, how to evaluate AI output quality) that users could apply outside of Harvey?
3. For firms using the API, is there a roadmap for publishing official SDKs or providing more granular control over model selection and retrieval parameters, so technical teams can build transferable AI engineering skills while working within the Harvey ecosystem?

EXIT

Portability & Migration

~ Partially Met

ASSESSMENT

Harvey exhibits a classic "roach motel" pattern: data goes in easily via standard formats but cannot be fully extracted. There is no bulk download capability for original source documents via UI or API—the platform lacks a documented content download endpoint, only offering metadata retrieval and file deletion. Review table

outputs can be exported to Excel, and query/usage history is available via API in JSON, which provides partial portability of work product. However, the most significant lock-in risk lies in custom workflows: the proprietary Workflow Engine stores all orchestration logic, prompt engineering, and "golden document" configurations internally with no export capability whatsoever. Organizations that invest heavily in building Harvey workflows will face a complete rebuild if they migrate, losing not just configuration but the iterative refinement embedded in those workflows. [80, 39, 81, 3, 1, 82, 42, 83, 18, 31, 33, 6, 31, 31, 35, 84, 85, 86, 6, 31, 3, 5, 3, 14, 9, 5, 13, 12, 7, 30, 5, 7, 11, 39, 5, 25, 87, 26, 5, 5, 16, 28]

TRADE-OFFS

● WHAT YOU GAIN

A purpose-built legal AI platform with proprietary workflow orchestration, DMS integrations, and domain-specific capabilities that would be difficult to replicate with general-purpose tools. The closed system likely contributes to the polished, low-friction user experience.

● WHAT YOU GIVE UP

Meaningful exit costs. Original documents stored in Vault have no programmatic bulk export path—you must rely on DMS sync as a workaround or negotiate a custom data extraction. All custom workflow logic, prompt engineering, and institutional process knowledge encoded in Harvey's Workflow Engine is completely non-portable. The longer and deeper you use the platform, the higher the switching cost becomes, which shifts negotiating leverage to the vendor at renewal.

QUESTIONS FOR VENDOR

1. Can you provide a contractual commitment to deliver a complete bulk export of all source documents, generated artifacts, and workflow configurations in standard formats (e.g., ZIP archive, structured JSON/YAML) within a defined timeframe upon contract termination?
2. For organizations that have built 50+ custom workflows on your platform, what is the concrete migration path? Can you export the underlying prompt logic, decision trees, and golden document references in any structured format—even a proprietary one we could parse?
3. If we use DMS integration (iManage/SharePoint) as the primary document store rather than uploading directly to Vault, does that fully mitigate the document extraction risk, or are there derived data assets (annotations, extracted metadata, review results) that would still be trapped in the platform?

Trade-off Summary

Criterion	Score	What You Gain	What You Give Up

SEE	~ A tightly controlled, legally-optimized system where Harvey's domain expertise manages the complexity of model selection, inference tuning, and retrieval—reducing the risk of misconfiguration in high-stakes legal work. The citation engine and Thinking States provide output-level verifiability without requiring users to understand AI internals. For firms that want reliable legal AI without building internal AI engineering capacity, this abstraction is a feature. The organization cannot independently audit, benchmark, or understand the AI mechanics driving its legal work product. If Harvey silently changes models, adjusts inference parameters, or modifies retrieval strategies, the firm has no way to detect or evaluate those changes. This also impedes the organization's ability to develop deep AI fluency—lawyers learn to use Harvey's interface, not to understand how AI-assisted legal reasoning actually works. The ~100 parallel model calls per prompt also raise unanswered questions about cost structure and whether Harvey has incentives to reduce that compute intensity over time in ways that degrade quality invisibly.
CHANGE	A fully managed, opaque AI layer eliminates the need for in-house ML engineering talent to operate the tool. Harvey's automated model orchestration ("100+ model calls" per query, mixture-of-experts routing) likely delivers strong out-of-the-box performance for standard legal tasks without any tuning burden. Legal professionals can use Playbooks and Workflows without technical expertise. The locked-down architecture also reduces the risk of misconfiguration by end users in a high-stakes legal context. The organization surrenders all control over AI behavior to Harvey. You cannot optimize for your firm's specific quality standards, risk tolerances, or cost sensitivities at the model level. If Harvey changes its underlying models, prompts, or retrieval logic, you have no visibility into or control over those changes. This creates deep dependency: your legal workflows are shaped entirely by Harvey's engineering decisions. It also impedes organizational AI learning—your teams build fluency in Harvey's interface, not in the underlying AI mechanics that will define competitive advantage. If Harvey's opaque orchestration produces suboptimal results for a specific use case, your only recourse is to rephrase your prompt or wait for Harvey to fix it.
ADAPT	~ A deeply integrated, enterprise-grade legal AI platform that rapidly incorporates OpenAI's latest capabilities (Deep Research, reasoning models) with legal-domain fine-tuning, modular workflow architecture, and a continuous release cadence—delivering frontier OpenAI features faster than most competitors can. Model optionality. If a non-OpenAI model becomes materially superior for legal reasoning, summarization, or any core task, Harvey cannot adopt it. The organization is making a second-order bet—not just on Harvey, but on OpenAI maintaining frontier leadership indefinitely. The fine-tuning sunk cost dynamic also creates perverse incentives for Harvey to stick with custom models even when newer base models outperform them.

USE	~	A purpose-built legal AI platform with enterprise-grade security, private document search (Vault/RAG), legal-specific fine-tuning that reduces hallucinations on case law, and pre-built workflows for high-value tasks like due diligence and contract review—all without requiring internal engineering to build and maintain a custom LLM stack.	Premium pricing for capabilities that partially overlap with commodity foundation models improving rapidly; risk that the differentiation gap narrows with each GPT/Claude update; potential adoption friction that limits realized ROI; and organizational dependency on Harvey's platform rather than building internal AI fluency that could be applied more broadly across the business.
LEARN	~	Lawyers and professional staff become productive with AI-assisted legal work quickly, without needing to understand underlying AI mechanics. The abstraction layer removes friction and reduces error rates for non-technical users, delivering immediate workflow value.	The organization builds proficiency in Harvey's interface rather than in AI itself. If Harvey is displaced or the firm wants to build internal AI capabilities, users will lack the foundational skills—prompt architecture, RAG design, model evaluation, context management—that would transfer. The firm becomes dependent on Harvey's engineering team for the AI competence that increasingly matters strategically.
EXIT	~	A purpose-built legal AI platform with proprietary workflow orchestration, DMS integrations, and domain-specific capabilities that would be difficult to replicate with general-purpose tools. The closed system likely contributes to the polished, low-friction user experience.	Meaningful exit costs. Original documents stored in Vault have no programmatic bulk export path—you must rely on DMS sync as a workaround or negotiate a custom data extraction. All custom workflow logic, prompt engineering, and institutional process knowledge encoded in Harvey's Workflow Engine is completely non-portable. The longer and deeper you use the platform, the higher the switching cost becomes, which shifts negotiating leverage to the vendor at renewal.

Key Questions for Your Decision

Consider these strategic questions based on your organizational context and priorities.

- 1 Is your firm's primary goal to deploy legal AI quickly and capture immediate productivity gains, or to build deep internal AI fluency that positions the firm competitively over the next 5-10 years—and how does a fully managed, opaque platform like Harvey align with that priority?
- 2 How comfortable are you making a second-order bet on OpenAI's continued frontier dominance, given that Harvey's exclusive partnership means you cannot access potentially superior models from Anthropic, Google, or open-source alternatives if the landscape shifts?

- 3 What is your firm's realistic capacity and appetite for building internal AI engineering talent—and if that capacity is low today, does outsourcing AI orchestration entirely to Harvey solve the problem or merely defer it while raising the eventual cost of catching up?
- 4 How would you detect and respond if Harvey silently degrades inference quality—through model swaps, reduced compute per query, or modified retrieval logic—given that the platform provides no visibility into its AI mechanics beyond output-level citations?
- 5 If Harvey were to be acquired, sunset a product line, or significantly raise prices at renewal, what is your concrete plan for extracting your institutional knowledge, custom workflows, and document corpus from the platform—and have you stress-tested that plan?
- 6 Does Harvey replace a capability gap that your firm genuinely cannot fill otherwise, or does it overlap significantly with what your lawyers could accomplish using rapidly improving general-purpose AI tools (e.g., ChatGPT, Copilot) combined with modest internal prompt engineering and RAG infrastructure?

References

- [1] <https://developers.harvey.ai/guides/assistant>
- [2] <https://developers.harvey.ai/api-reference/authentication>
- [3] <https://developers.harvey.ai/api-reference/query-vault-knowledge-source/completion-with-vault-rag>
- [4] <https://help.harvey.ai/articles/harvey-guide>
- [5] <https://help.harvey.ai/articles/harvey-guide>
- [6] <https://developers.harvey.ai/api-reference/completion/completion>
- [7] https://www.harvey.ai/_next/image
- [8] <https://images.ctfassets.net/kftzwdyauwt9/6t4FNwBfJZjReOB011Rxot/e04d453fe93d075b482edf196bfb5110/harvey.webp>
- [9] <https://www.harvey.ai/blog/how-in-house-legal-teams-build-the-case-for-ai-adoption>
- [10] <https://www.cnbc.com/quotes/GOOGL>
- [11] <https://www.harvey.ai/platform/workflows>
- [12] <https://www.harvey.ai/blog/integrating-deep-research-into-harvey>
- [13] <https://www.harvey.ai/blog/rebuilding-harveys-design-system-from-the-ground-up>
- [14] <https://www.harvey.ai>
- [15] https://en.wikipedia.org/wiki/Google_DeepMind
- [16] <https://www.harvey.ai/resources/videos>
- [17] https://images.ctfassets.net/kftzwdyauwt9/629pWHy9yqhzLk6TS7elo8/79af6b547ddf6f2c265cd7739bf9b3cf/Harvey_SBS.gif
- [18] <https://api.harvey.ai/api/v2/completion>
- [19] <https://openai.com/blog/introducing-improvements-to-the-fine-tuning-api-and-expanding-our-custom-models-program>
- [20] <https://preview.redd.it/sv5bjrtuxqhg1.jpeg>
- [21] <https://www.reddit.com/r/legaltech/comments/1ku1gh8/comment/mw7uvag>
- [22] <https://www.harvey.ai/blog/biglaw-bench-sources>
- [23] <https://www.reddit.com/r/legaltech/comments/1ku1gh8/comment/mu225g9>
- [24] <https://www.harvey.ai/solutions/innovation>
- [25] <https://www.harvey.ai/platform/ecosystem>
- [26] <https://help.harvey.ai/topics/getting-started>
- [27] <https://help.harvey.ai>

- [28] <https://www.harvey.ai/enterprise>
- [29] <https://help.harvey.ai/roadmap>
- [30] <https://www.harvey.ai/platform/assistant>
- [31] <https://developers.harvey.ai/guides/vault>
- [32] <https://developers.harvey.ai/api-reference/retrieve/retrieve-latest-audit-log>
- [33] <https://developers.harvey.ai/api-reference/vault/delete-an-entire-vault-project-and-all-its-contents>
- [34] <https://openai.com/index/harvey>
- [35] <https://developers.harvey.ai/api-reference/vault/list-vault-projects-and-knowledge-bases>
- [36] https://developers.harvey.ai/guides/client_matters
- [37] <https://www.cnbc.com/2025/08/04/legal-ai-startup-harvey-revenue.html>
- [38] <https://www.reddit.com/user/Librality>
- [39] https://developers.harvey.ai/guides/usage_history
- [40] <https://developers.harvey.ai/api-reference/retrieve/retrieve-earliest-audit-log>
- [41] https://developers.harvey.ai/guides/audit_logs
- [42] https://mintcdn.com/counselai/4kvXwgpYqzpV1Xq9/logo/Harvey_Logo_Ink.svg
- [43] <https://www.clio.com/blog/legal-innovation>
- [44] <https://openai.com/contact-sales>
- [45] https://styles.redditmedia.com/t5_2m15zd/styles/profileIcon_4xp901rq1e651.jpg
- [46] <https://www.reddit.com/r/legaltech/comments/1ku1gh8/comment/mu1yoo1>
- [47] <https://novehiclesinthepark.substack.com>
- [48] <https://www.reddit.com/r/legaltech/comments/1bvtuic/comment/ky1xifj>
- [49] <https://www.reddit.com/r/legaltech/comments/1ku1gh8/comment/mu2lmzn>
- [50] <https://www.harvey.ai/blog/mobile-voice-announcement>
- [51] <https://www.reddit.com/r/legaltech/comments/1ku1gh8/comment/mu03fjl>
- [52] https://styles.redditmedia.com/t5_2079d2/styles/profileIcon_qzuwhfh5lsdc1.jpg
- [53] <https://www.reddit.com/search>
- [54] <https://www.reddit.com/r/legaltech/comments/1ku1gh8/comment/mtzs8xd>
- [55] <https://www.clio.com/features/legal-ai-software>
- [56] <https://www.harvey.ai/legal>
- [57] <https://developers.harvey.ai/guides/introduction>
- [58] [https://en.wikipedia.org/wiki/Harvey_\(software\)](https://en.wikipedia.org/wiki/Harvey_(software))
- [59] https://www.reddit.com/r/legaltech/comments/1ku1gh8/harvey_ai_reviews_general_advice_for_a
- [60] https://en.wikipedia.org/wiki/Ashurst_LLP
- [61] <https://www.reddit.com/user/alexdenne>
- [62] <https://developers.harvey.ai/api-reference/usage/get-usage-history>
- [63] https://styles.redditmedia.com/t5_bak5sq/styles/profileIcon_e0kmysttkghf1.png
- [64] <https://www.harvey.ai/customers/the-adecco-group>
- [65] <https://cdn.sanity.io/images/07s0r5r6/production/36bbd399121eff2b23cd36789ee677a065c3bd68-284x200.svg>
- [66] <https://www.harvey.ai/company/careers/15dfbe0b-f6e2-4bee-a906-65e31dfc5b23>
- [67] <https://www.harvey.ai/blog/how-harvey-saves-lawyers-time>
- [68] <https://www.harvey.ai/solutions/transactional>
- [69] <https://alb.reddit.com/cr>
- [70] https://styles.redditmedia.com/t5_8uzy8o/styles/profileIcon_8z5d1hbl2jwd1.png
- [71] <https://www.clio.com/wp-content/uploads/2025/06/CoCounsel-ThomsonReuters-Featured-e1750802289894-750x387.png>
- [72] <https://www.reddit.com/r/legaltech/comments/1ku1gh8/comment/mu1tf3h>
- [73] <https://www.clio.com/resources/ai-for-lawyers/chat-gpt>
- [74] <https://www.clio.com/blog/harvey-ai-legal>
- [75] <https://www.clio.com/enterprise/vincent>
- [76] https://help.harvey.ai/_next/image
- [77] <https://help.harvey.ai/articles/getting-started-with-harvey>
- [78] https://www.reddit.com/user/D_Lex

- [79] <https://www.reddit.com/r/legaltech/comments/1ku1gh8/comment/mu3gbwv>
- [80] <https://developers.harvey.ai/api-reference/ratelimits>
- [81] <https://www.harvey.ai/help>
- [82] <https://developers.harvey.ai/api-reference/vault/upload-files-to-project>
- [83] https://api.harvey.ai/api/v1/vault/get_row/%7Breview_event_id%7D/%7Bfile_id%7D
- [84] <https://api.harvey.ai/api/v1/vault/workspace/projects>
- [85] <https://developers.harvey.ai/api-reference/add/add-client-matters>
- [86] <https://developers.harvey.ai/api-reference/vault/create-a-new-vault-project>
- [87] <https://www.youtube.com/watch>

CONFIDENTIAL — FOR INTERNAL USE ONLY

This evaluation framework assesses vendors on immediate utility and long-term organizational AI capacity building.

Framework designed for Fortune 500 executive decision-making • Version 1.0